

TOWARDS AN ENERGY EFFICIENT COMPUTATION IN ERLANG

Áron Attila Mészáros

Gergely Nagy



ELTE IK, Computer Science BSc

STCS 2019


MOTIVATION

- Environmental awareness
 - Energy conservation, minimizing energy consumption
 - Even in computers
 - Green computing
- Why Erlang?
 - Popular languages – lots of research (e.g. C++, Java, Haskell)
 - Erlang
 - Widely used in the industry
 - No such research yet



PREVIOUS WORK

- Tool for measuring energy consumption
- Measuring complex algorithms
 - N-queens
 - Sparse matrix multiplication
- Special attention to:
 - Lists vs arrays
 - Higher order functions
 - Parallelization (parallel map)



$$\begin{bmatrix} 0 & 0 & 0 & 0 \\ 5 & 8 & 0 & 0 \\ 0 & 0 & 3 & 0 \\ 0 & 6 & 0 & 0 \end{bmatrix} \times \begin{bmatrix} 3 & 0 & 0 & 0 \\ 0 & 0 & 2 & 0 \\ 5 & 3 & 0 & 7 \\ 0 & 1 & 0 & 0 \end{bmatrix} = \begin{bmatrix} 0 & 0 & 0 & 0 \\ 15 & 0 & 16 & 0 \\ 15 & 9 & 0 & 21 \\ 0 & 0 & 12 & 0 \end{bmatrix}$$

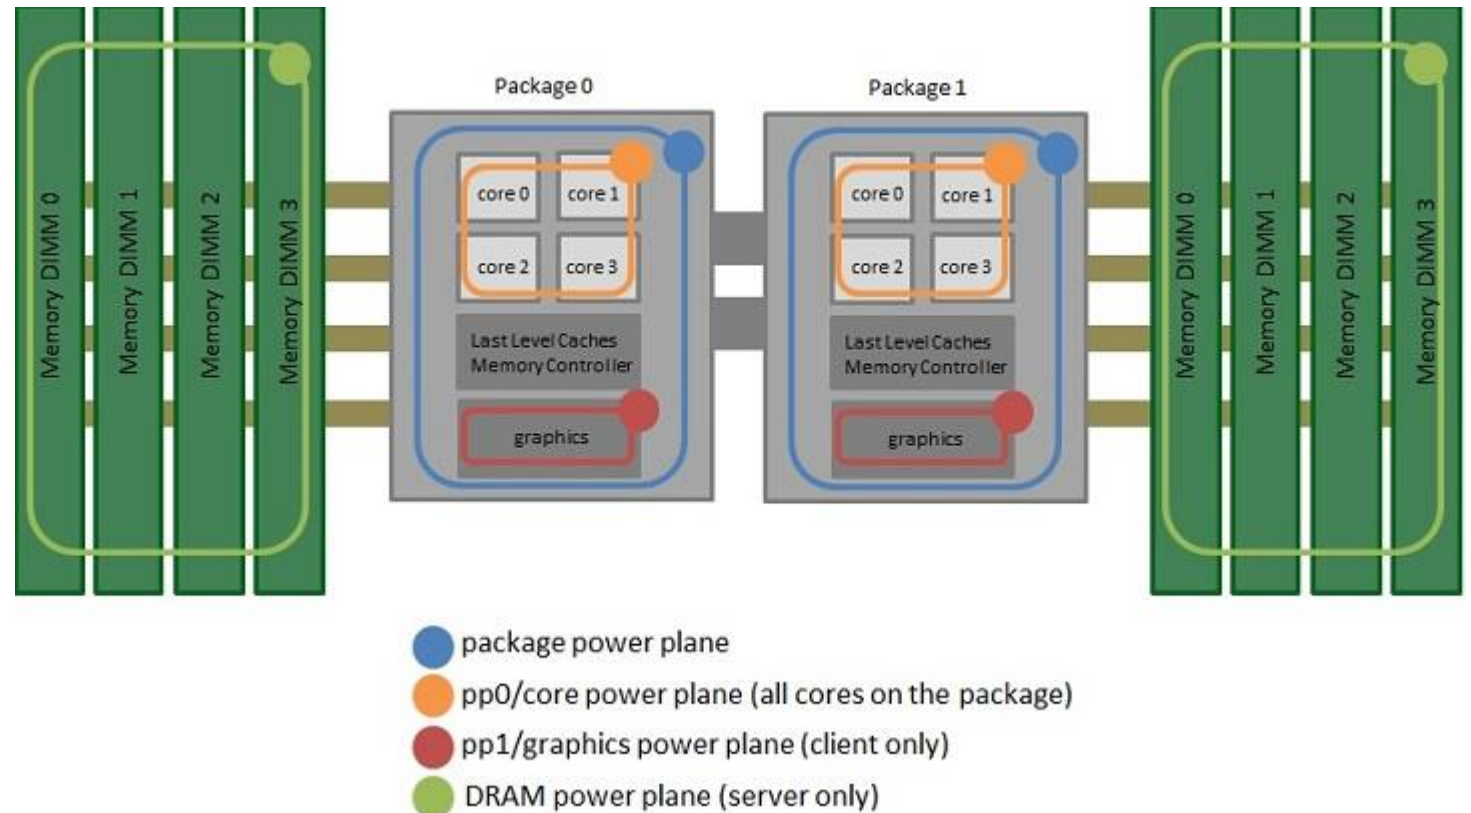
NEW GOALS

- Improve methodology
- Directly measure language elements
- Measure data structures independently
- Further inspection of parallelization
 - Spawning processes
 - Sending messages
- Refactoring in order to minimize energy consumption

HOW TO MEASURE ENERGY CONSUMPTION

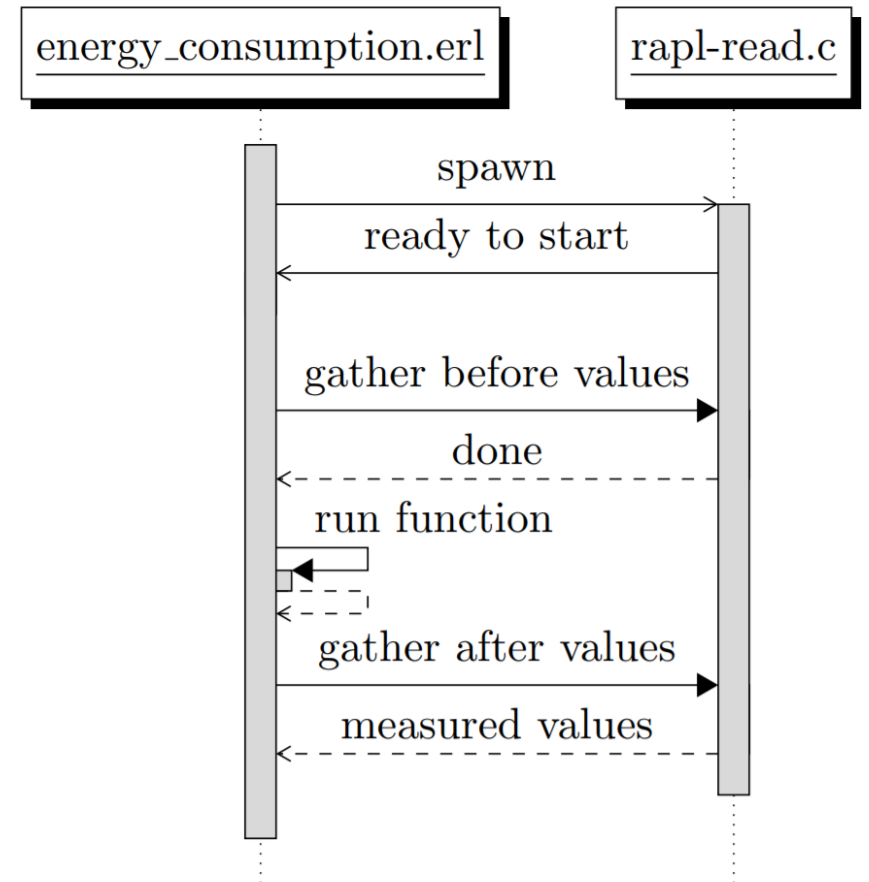
- Running Average Power Limit (RAPL)

- Intel
- Methods:
 - MSR
 - perf_event,
 - sysfs
- Domains:
 - PKG – package
 - PP0 – core
 - PPI – uncore
 - DRAM



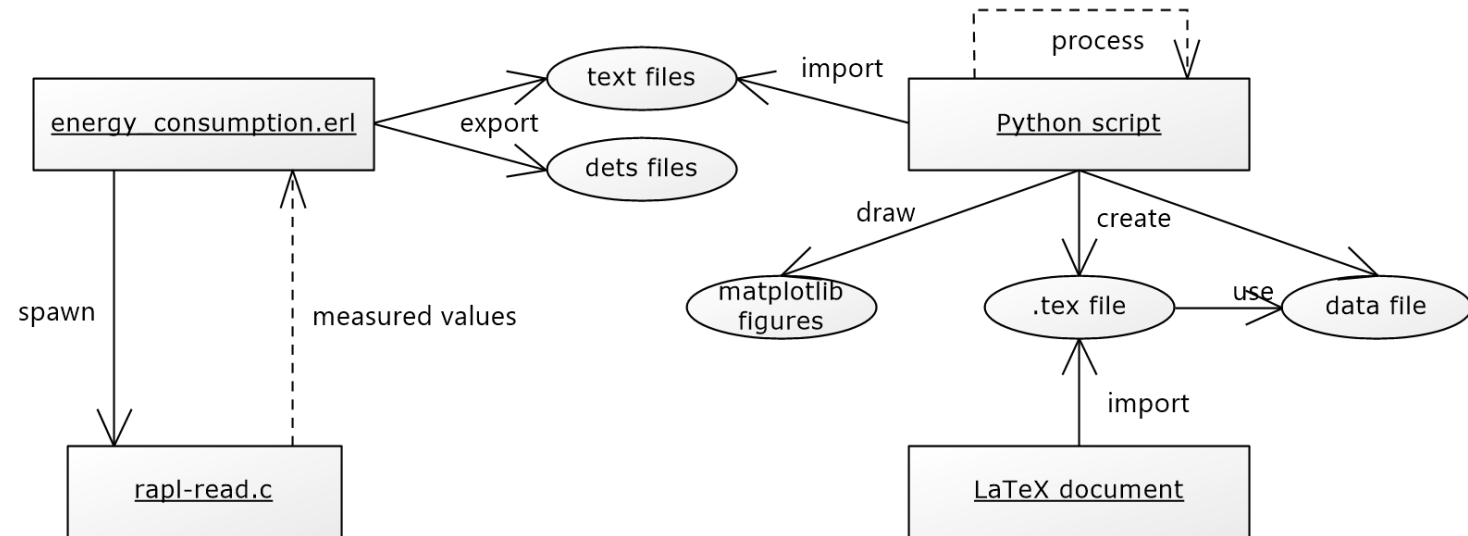
ERLANG FRAMEWORK

- C program to read RAPL values
- Erlang module to run measured functions
 - Communication via ports
 - Runs the functions
 - Tells the C program when to measure
- Python GUI
 - Using TkInter
 - Setup measurement
 - Helps with organizing the measurements
 - Visualizing results



PROCESSING DATA

- Methodology
 - Everything is measured 10 times
 - Discarding minimal and maximal values
 - Calculate average energy consumed
 - Also measuring runtime
- Visualization
 - Python – matplotlib



HIGHER ORDER FUNCTIONS

- Parameters or the result is a function
- Integral part of functional programming languages
- Examples
 - map
 - filter
 - fold
- We wanted to find the cost of passing functions as a parameter

HOFs - MEASUREMENTS

- Passing named or unnamed function
- Using HOF or list comprehension or recursion
- Implementing our own higher order functions
- Measured:


- map
- filter
- map ◦ filter

```
lists:map(fun increase/1, L).
```

```
lists:filter(fun even/1, L).
```


```
lists:map(fun increase/1, lists:filter(fun even/1, L)).
```

MAP - RESULTS




- Best:
 - List comprehension
 - Recursion
- Worst:
 - List comprehension with unnamed function
 - All map implementations
- Ranking is same based on energy and time

MAP - RESULTS

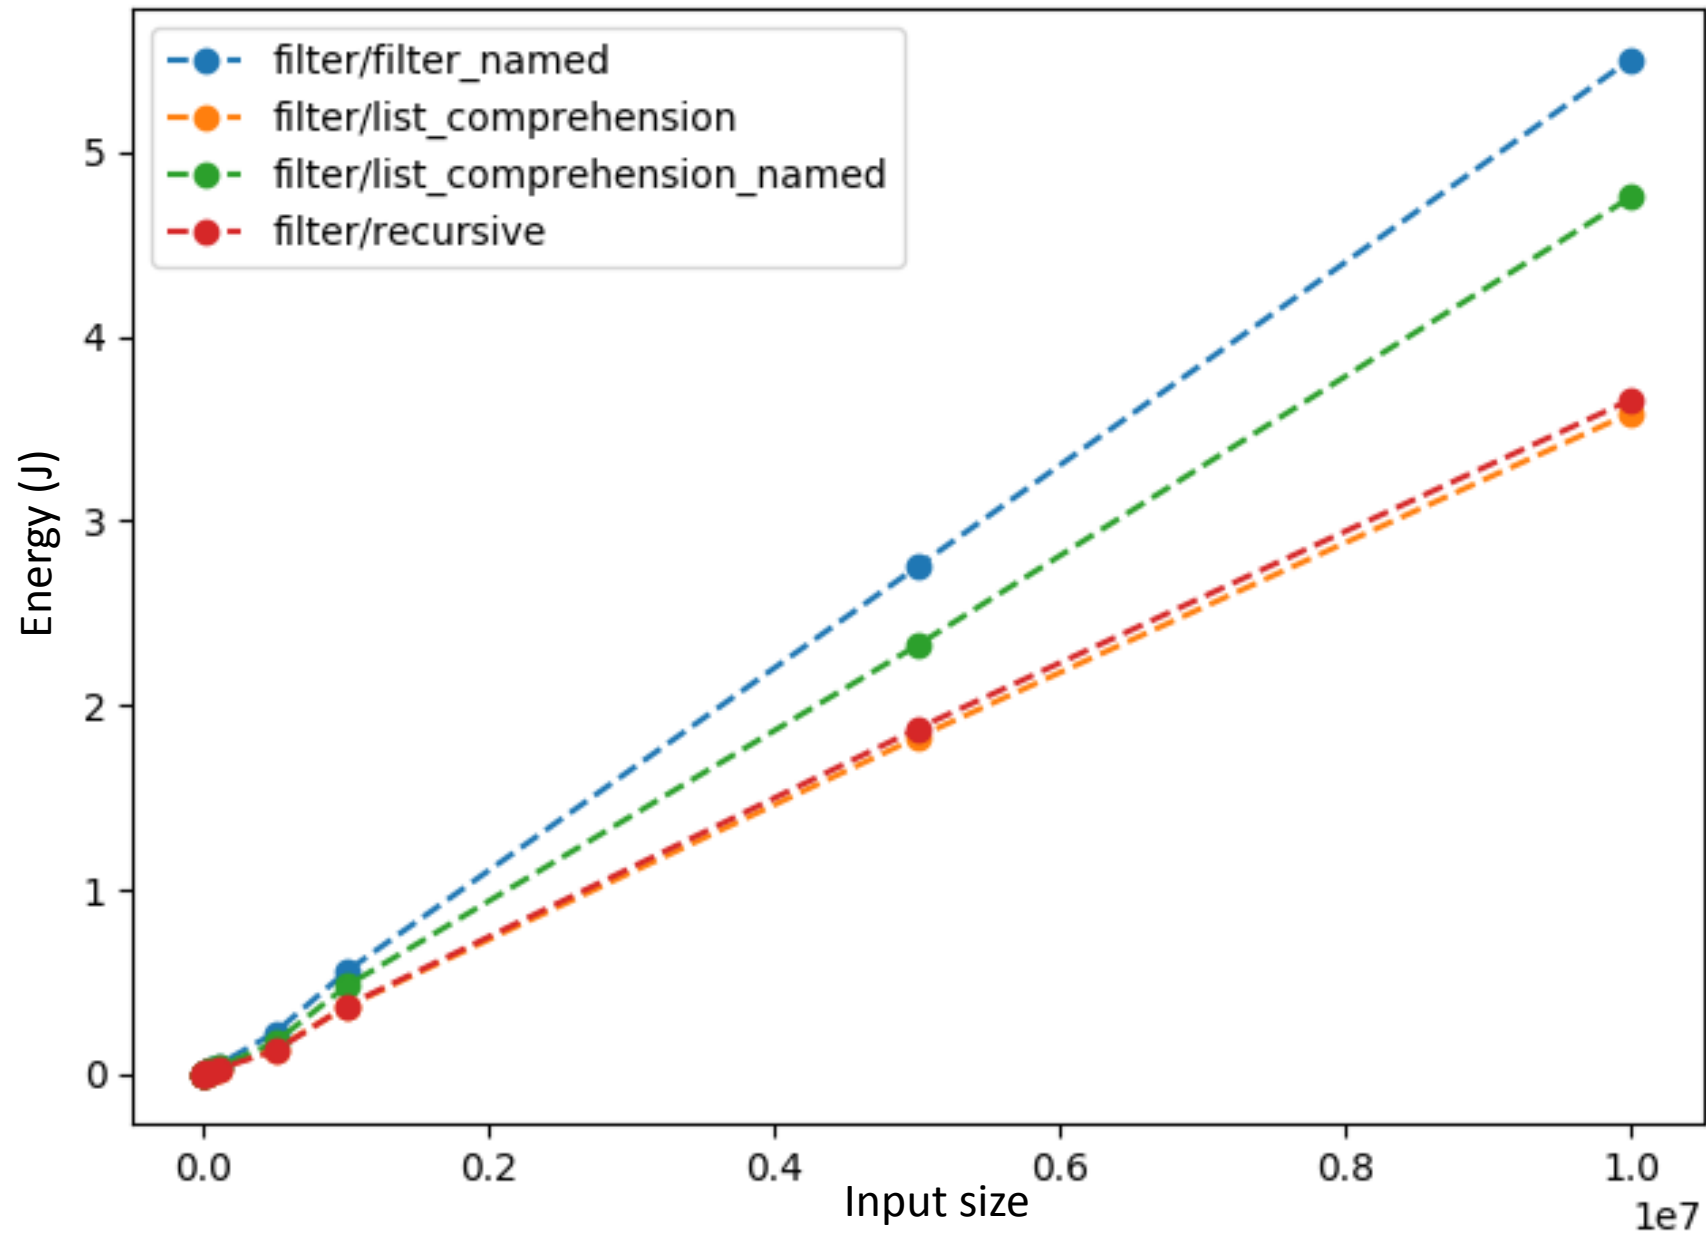


FILTER - RESULTS

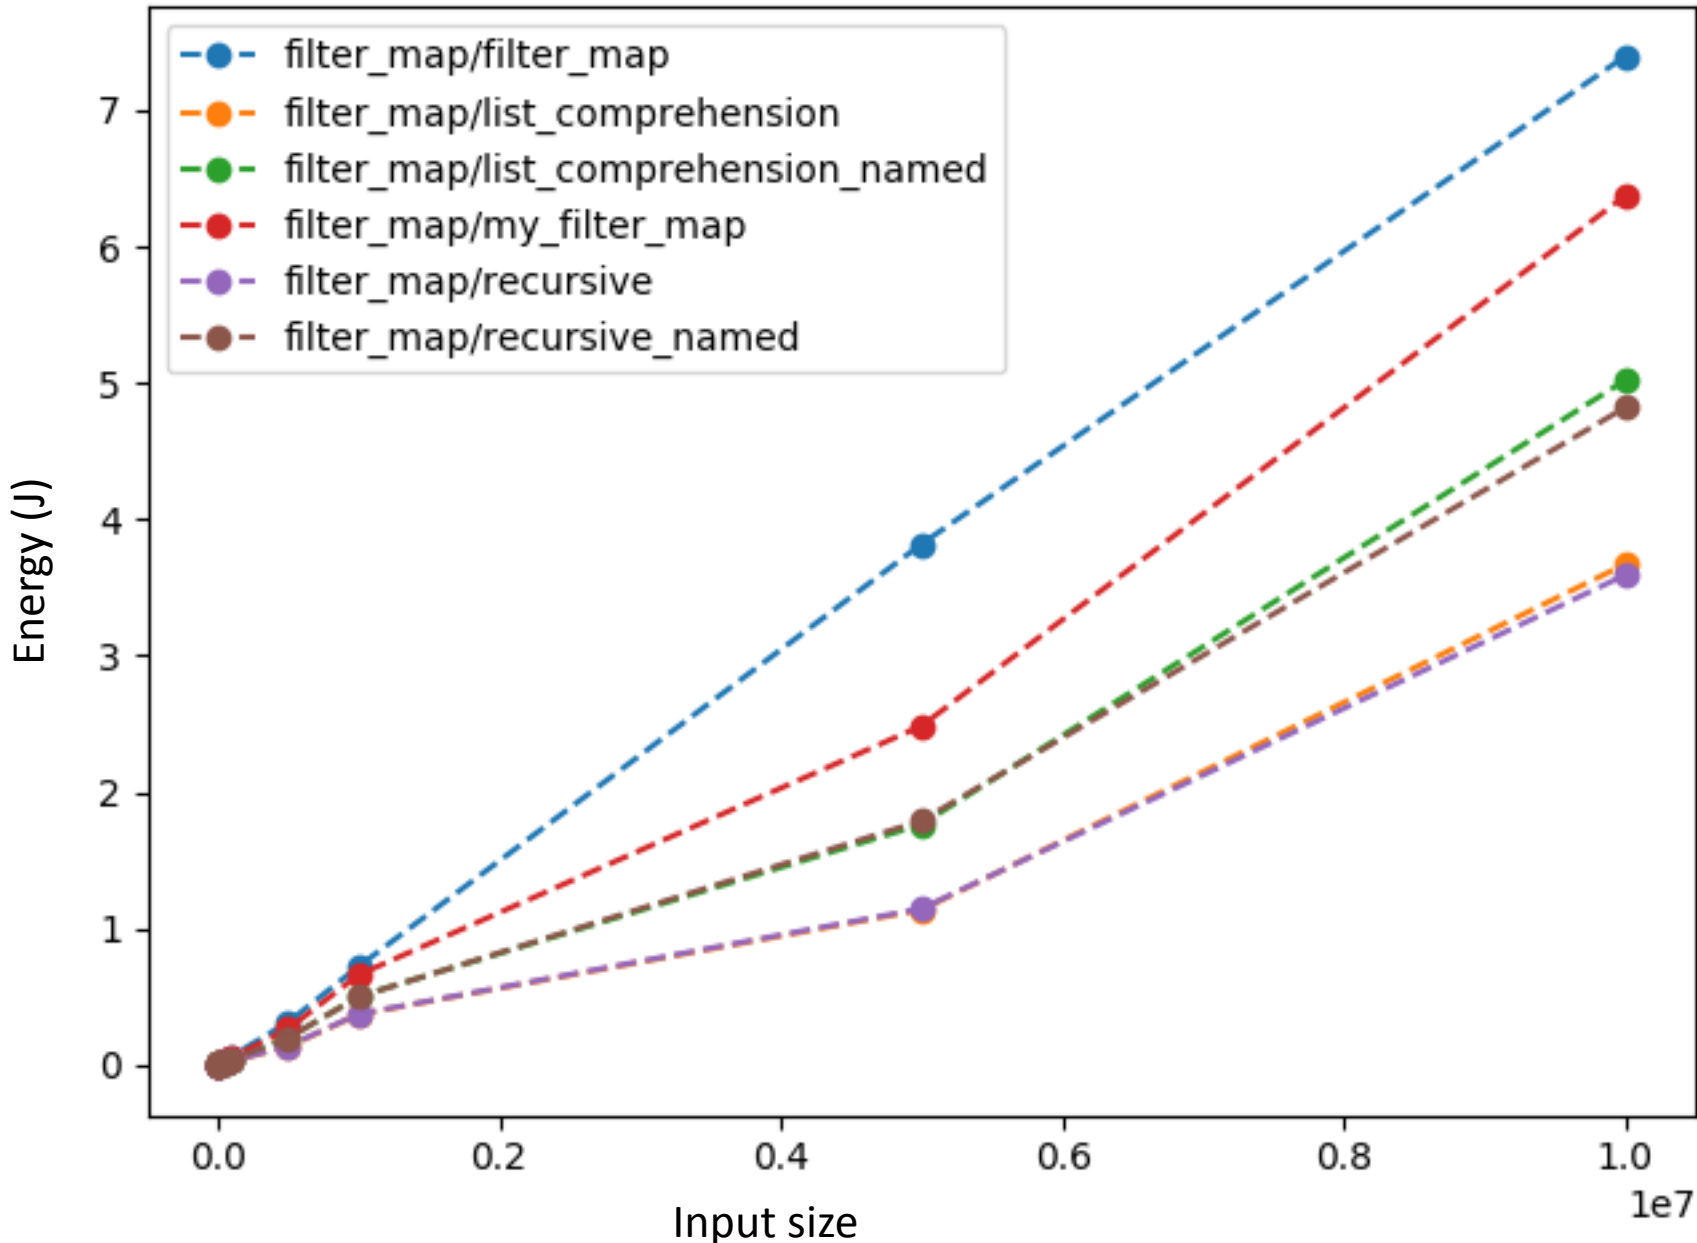


- Best:
 - List comprehension
 - Recursion
- Worst:
 - List comprehension with unnamed function
 - All map implementations
- Ranking is same based on energy and time
- Same trends as with map

FILTER - RESULTS



MAP ◦ FILTER - RESULTS



- Best:
 - List comprehension
 - Recursion
- Worst:
 - Filter then map
 - Own HOF implementation
- Ranking is same based on energy and time


HOFs - FINDINGS

- Refactoring higher order function calls to a simple recursive function or list comprehension reduces energy consumption and runtime
- Eliminating the HOF is not enough, we need to eliminate the called function as well to minimize energy consumption.
- This is consistent with previous findings (N-queens)

DATA STRUCTURES



- Investigating different ways to store key-value pairs
 - list of tuples
 - map
 - dict
- Measuring different operations on those data structures
 - find
 - delete
 - update
 - building the data structure

BUILDING - RESULTS




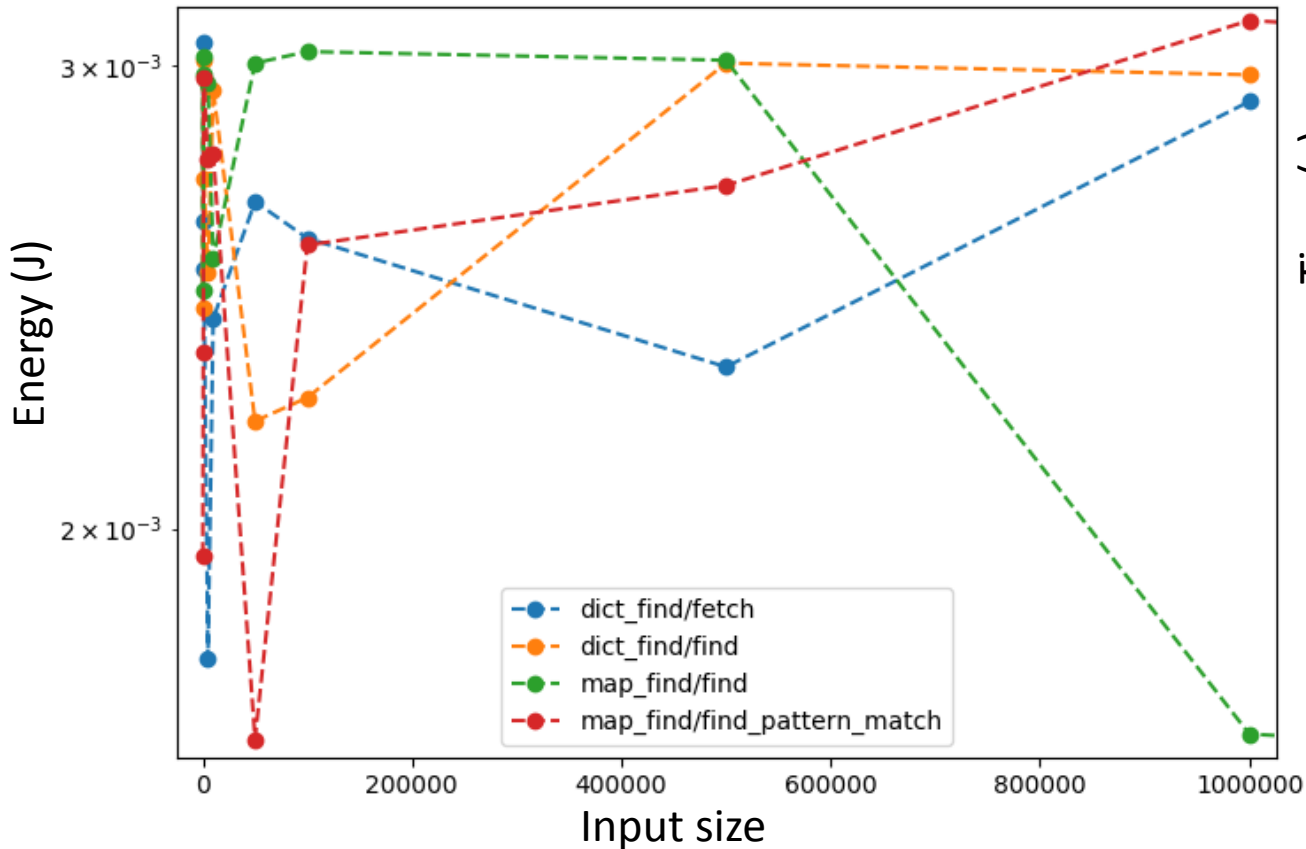
- Dicts are really expensive to build
- Lists are the cheapest
- Building a map from a list of tuples is cheaper than building the map directly

FIND - RESULTS




- Finding with pattern matching is faster, but consumes more energy
- Using the find function consumes the least energy but is slower

FIND - RESULTS





- Finding in a dict is similar to finding in a map,
- Maps are quicker
- Energy consumption may be better for maps (cannot construct large enough dicts)

FIND - RESULTS



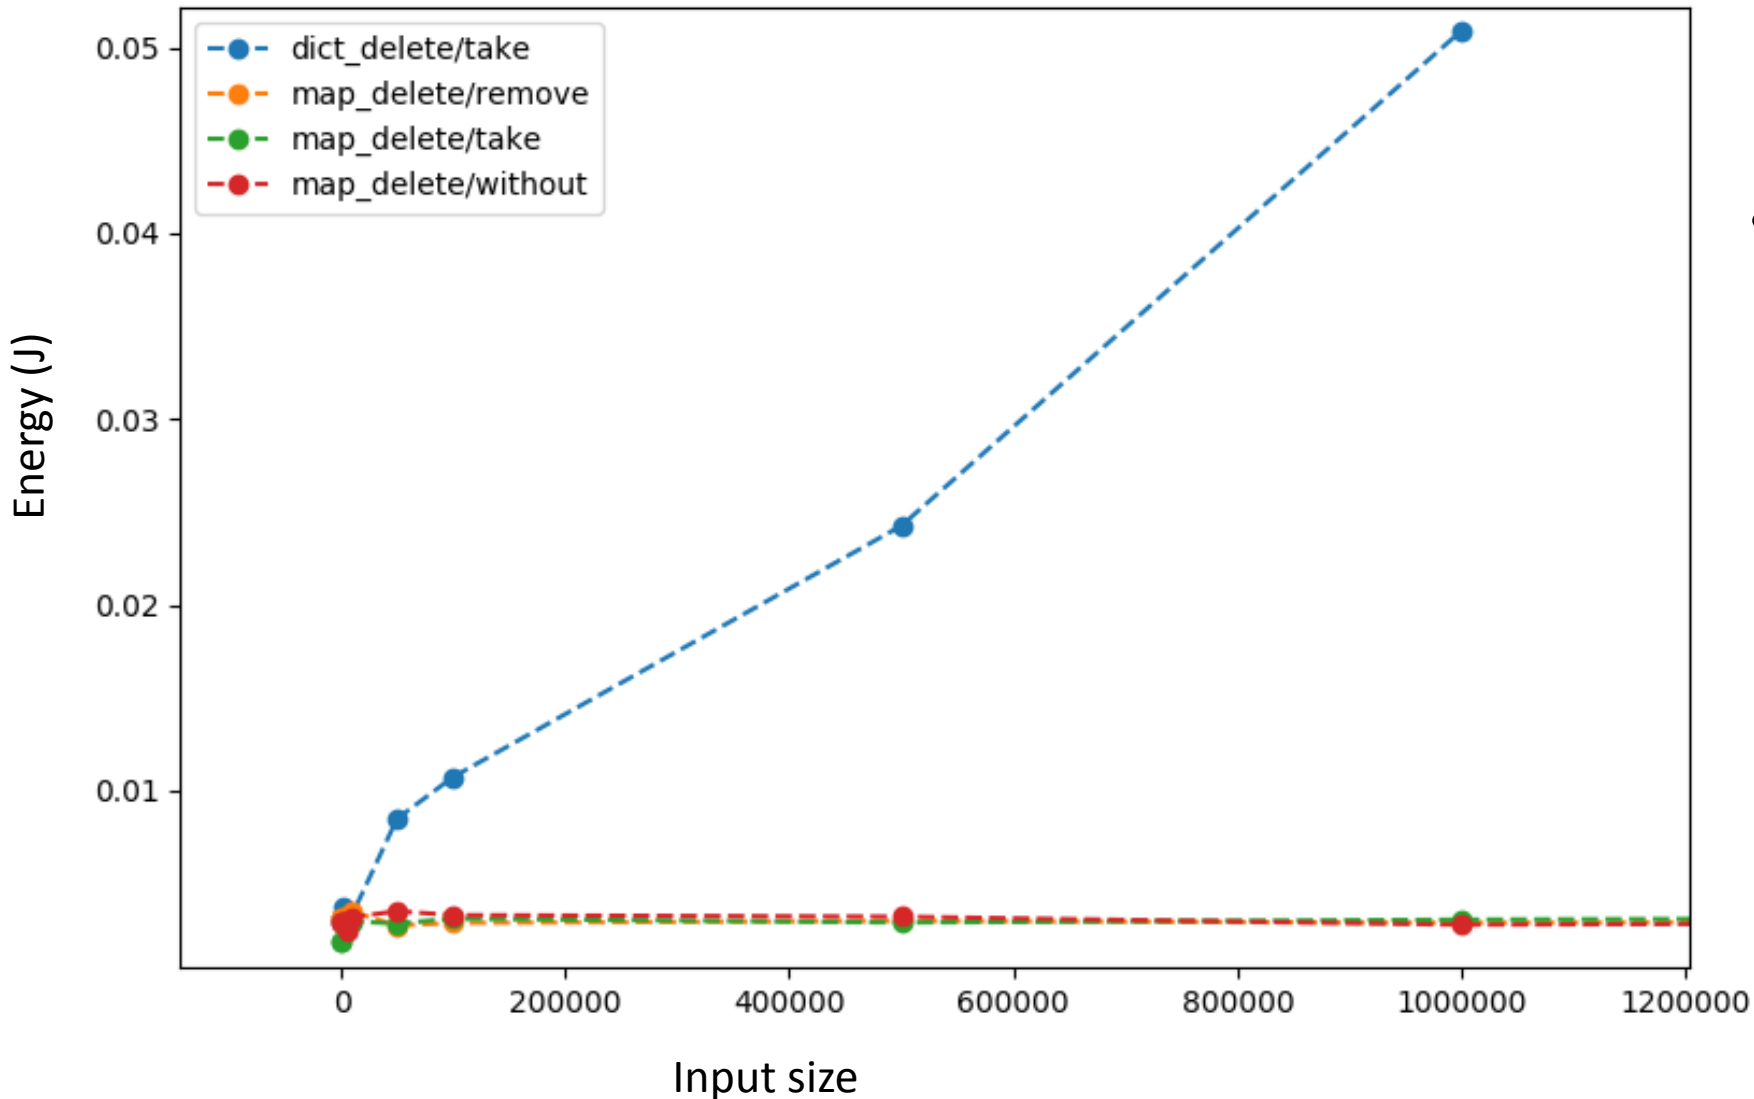
- proplists:get_value is inefficient, even a simple recursive find is better
- lists:keyfind is better
- Compared to maps all are worse (expected, because of complexity differences)

DELETE - RESULTS




- Deleting with `maps:without` is the slowest, but consumes the least energy

DELETE - RESULTS



- Deleting from a dict is worse than deleting from a map
- Complexity difference

DELETE - RESULTS



- Deleting from list is much worse than all previous examples
- Complexity difference
- No conclusive difference between the two ways to delete

UPDATE - RESULTS

- Results for updating an element in our dictionaries show the same patterns as deleting an element.
- Maps are the most efficient
- Dicts are in the middle
- Lists perform the worst

DATA STRUCTURES - FINDINGS

- Most of the differences between the data structures can be attributed to different representations
- Interesting findings are:
 - We should use `lists:keyfind` instead of `proplists:get_value`
 - With maps the quickest solution is not always the more efficient
 - More investigation is needed, because measured values are so small
- Refactoring may be possible between different implementations of functions on the same data structure
- Based on the cost to build every data structure we may be able to decide how many of each operation is needed until it is worth building a different data structure

CONCLUSIONS

- Tool for measuring energy consumption
- Higher order functions have an overhead
 - It may be eliminated by rewriting the calls to higher order functions with recursion or list comprehensions
 - Function calls should also be eliminated
- Erlang built-in functions on the same data structures can have significant differences in performance

FUTURE WORK

- Investigating the cost of spawning processes and sending messages
- Measuring different algorithmic skeletons for parallelization
- Creating refactorings based on our findings using RefactorErl

THANK YOU FOR YOUR ATTENTION